

Pair a compatible AI assistant with your project to build and edit your schedule in real time · scheduler.spertsuite.com

WHAT IS CONNECT AI?

Connect AI is a feature in SPERT® Scheduler that lets a compatible AI assistant build or edit your project schedule in real time. When the connection is active, the AI can create activities with a three-point estimate (min, most likely, max — in working days) and a recommended probability distribution, rename activities and adjust estimates, write or update a short description of what an activity entails, add checklist and deliverable items, append notes explaining its reasoning, create milestones and assign activities to them, reorder a scenario's activities, and — once you turn on dependency-aware scheduling for the scenario — link activities together with dependencies. For a large schedule, the AI can also work in **bulk** — creating or updating many activities, milestones, assignments, or dependencies in a single call, appending a note to many existing activities at once, or importing a whole schedule at once. Every change the AI makes appears in your project as it happens, alongside a live feed showing exactly what was done. You stay in control at every step: the assistant asks what you want to build before doing anything, and you choose whether it can also read your current project.

Connect AI works through the **MCP protocol** (Model Context Protocol) — a standard for giving AI assistants access to external tools. SPERT Scheduler shares a single MCP server with the rest of the SPERT Suite, at mcp.spertsuite.com, that the AI calls to read and write your schedule. The AI assistant must be configured to reach this server before the feature can be used — the same one-time setup works for every SPERT app.

HOW IT WORKS — EACH SESSION

#	Action	Detail
1	Open a project and click Connect AI	The Connect AI button is in the project header. A consent modal appears.
2	Choose permissions	Write Mode (required) — the AI can create and edit activities, estimates, milestones, and notes in this project. Read Mode (optional) — lets the AI read the current schedule and activity ids so it can make context-aware edits. Required for dependency edits and for reordering activities, and generally needed whenever the AI needs to reference an activity or milestone you (rather than it) created. Click Connect .
3	Copy the pairing code and prompt	A short pairing code appears alongside two copy buttons: Copy Code (the code alone) and Copy Prompt (a ready-to-paste message that includes the code and detailed instructions for the AI). The code expires in 15 minutes.
4	Paste the prompt into your AI assistant	Open your AI assistant — it must already be configured with the SPERT MCP server (see One-Time Setup). Paste the copied prompt and send it. The assistant will connect, confirm your project is open, then ask what you're scheduling and whether the activities have real dependencies between them before creating anything — it will not build until you answer.
5	Review and disconnect	Activities, estimates, milestones, notes, and dependencies appear in your project in real time, and each change is logged in the live activity feed as it happens. Use Ctrl+Z to undo — a run of AI changes with no edit of your own in between undoes as one step. Toggle Read Mode at any time from the Connect AI panel. Click Disconnect when done — or the session expires automatically after 7 days.

Note: Do not switch projects while the AI is actively working — Connect AI is scoped to the project that was open when you clicked Connect. Switching scenarios mid-session is fine for most changes, since the AI's edits normally land in whichever scenario is on screen when they arrive — but any dependency-linking request always names its own scenario explicitly, so it isn't affected by which one you're looking at.

Want to practice first? On the Projects page, **Load Sample Project** builds a complete worked example — a fourteen-month cloud ERP implementation with 40 activities, 68 dependencies, 4 milestones, and dependency-aware scheduling already switched on. It is an ordinary project, so you can point the AI at it and try anything in this guide without touching real work: turn on Read Mode and ask the AI to summarize the critical path, reorder the activities, or add notes across a whole phase. Delete it when you're done, or load a fresh copy any time — each one is fully independent.

ONE-TIME SETUP: CONNECT YOUR AI CLIENT

Connect AI requires an AI client that supports MCP servers. Complete this setup once; every subsequent session — in Scheduler or any other SPERT app — starts at step 1 above. Seven clients work today. Scheduler's bulk tools mean a large schedule usually takes a handful of calls rather than one per activity, so most clients hold up well even on a big build — **ChatGPT and Claude tend to hold up best** on the rare build that still needs a long run of individual calls (a schedule the AI has to revise item-by-item, for example); any connected assistant works well for a handful of activities or a quick edit.

claude.ai web (Pro / Max / Team / Enterprise)	Lowest-friction path for non-developers — no terminal or config file required. Free-tier accounts are not supported. In claude.ai, go to Settings → Connectors → Add custom connector . Paste the server URL: https://mcp.spertsuite.com No API key is needed — click Save. To use it: open a new conversation, confirm the SPERT connector is active, open SPERT Scheduler in your browser, click Connect AI, and paste the Copy Prompt.
Claude Code (CLI, any plan)	Best for developers — one command. Run once in any terminal: <code>claude mcp add --transport http spert-suite https://mcp.spertsuite.com</code> Add <code>--scope user</code> to make the server available across all projects (default is project-scope only). Confirm with <code>claude mcp list</code> — you should see spert-suite in the output. To use it: start a Claude Code session, open SPERT Scheduler, click Connect AI, and paste the Copy Prompt.
Claude Desktop (config file)	Requires Node.js (for npx). Locate the config file: on macOS, <code>~/Library/Application Support/Claude/claude_desktop_config.json</code> ; on Windows, <code>%APPDATA%\Claude\claude_desktop_config.json</code> . Add the following to the mcpServers object (create the file if it does not exist): <code>"spert-suite": { "command": "npx", "args": ["mcp-remote", "https://mcp.spertsuite.com"] }</code> Save the file and restart Claude Desktop. mcp-remote installs automatically on first run.
Grok (paid tiers)	Self-service on all paid Grok tiers — similar UX to claude.ai Connectors. Go to grok.com/connectors → New Connector → Custom → paste https://mcp.spertsuite.com . No authentication required — click Create. To use it: open a Grok conversation, open SPERT Scheduler in your browser, click Connect AI, and paste the Copy Prompt.
ChatGPT (Pro / Team / Enterprise / Edu, Developer Mode)	Requires Developer Mode to be enabled. On Pro plans, go to Settings → Apps & Connectors → Advanced → Developer Mode → On . On Business / Enterprise / Edu plans, a workspace admin must enable it from Workspace Settings → Permissions & Roles → Connected Data. Once enabled, add the server: click + in a conversation → More → paste https://mcp.spertsuite.com , no authentication. To use it: open SPERT Scheduler in your browser, click Connect AI, and paste the Copy Prompt into the ChatGPT conversation. Recommended when a build still needs many individual tool calls despite the bulk tools (e.g. a long, item-by-item revision pass) — it holds up well across a long run of sequential calls. Note: if ChatGPT shows fewer tools than expected or behaves unexpectedly, check that you are on a capable model tier; free-tier model downgrades can cause tool enumeration issues.

Copilot Studio (Microsoft 365, advanced)

In a Copilot Studio agent, go to **Tools** → **New tool** → **Model Context Protocol**, paste <https://mcp.spertsuite.com>, set Authentication to **None**, and click Create. Use the agent's Test panel to interact. Requires a Microsoft 365 license. **Note on Microsoft 365 Copilot Chat (the standard chat experience):** Copilot Chat connects and works well for a small number of activities, a quick edit, or a single bulk call, but may pause for repeated confirmation or stall on a build that still needs a long run of individual calls. For that case, use ChatGPT or Claude instead. Copilot Studio (this row) is a stronger orchestrator than Copilot Chat and handles longer sequences; it requires additional admin configuration to deploy to Copilot Chat.

Gemini CLI (developer path)

Google's command-line AI tool supports MCP via a config file, similar to Claude Code. Add the SPERT server to your Gemini CLI settings file under the **mcpServers** key with the server URL <https://mcp.spertsuite.com>. To use it: start a Gemini CLI session, open SPERT Scheduler, click Connect AI, and paste the Copy Prompt.

Choosing the right assistant for the job: most builds now fit in a handful of bulk calls, so most connected assistants hold up well. **ChatGPT and Claude** additionally hold up well across a long run of individual tool calls, for the less common build that still needs one — an item-by-item revision pass, for example. **Microsoft 365 Copilot Chat** is great for a quick, targeted edit or a single bulk call — one activity, one estimate, one import — but may stall or ask for repeated confirmation over a long run of individual calls. For that case, switch to ChatGPT or Claude. **Grok and Copilot Studio** connect and work with both the bulk tools and individual calls; behavior over a long sequential build is best confirmed with a short test run on your own account.

WHAT STILL DOESN'T WORK

Connect AI requires an AI client with MCP server support configured. If you paste the Copy Prompt into an AI that has no MCP connection, the assistant reads the tool-call instructions but has no actual channel to your project — nothing changes in the app. In some cases the AI may even fabricate a response, describing a schedule it claims to have built. The tell: tool names like **resolve_session_code** appear in the reply, but your project stays empty.

Not supported today: Gemini web app (gemini.google.com) — no self-service custom MCP connector in the consumer interface yet (Gemini CLI works; the web app does not). Free-tier claude.ai accounts — Connectors require a paid plan. ChatGPT Free and Plus without Developer Mode enabled. Any other AI assistant without MCP support configured.

WHAT'S COMING

MCP adoption has moved quickly. Seven clients work today: Claude (claude.ai, Code, Desktop), ChatGPT, Grok, Copilot Studio, and Gemini CLI. The remaining gap for consumer users is the Gemini web app — there is no self-service custom MCP connector in the consumer Gemini web interface yet, though Gemini CLI works for developers via a config file.

As MCP adoption continues to expand, Connect AI will work with new clients automatically — no changes to the SPERT server or your workflow will be required. The same pairing-code workflow works across all supported clients today, and across every SPERT app that offers Connect AI.

WRITE MODE, READ MODE & PRIVACY

Write Mode (required)

Lets the AI create and edit activities, estimates, milestones, descriptions, and notes in this project: create an activity with a three-point estimate and a recommended probability distribution; rename an activity or adjust its estimate; set or replace an activity's short plain-language description; append a note explaining its reasoning; add checklist or deliverable items and mark existing ones complete; create a milestone and assign or unassign activities to it; reorder a scenario's activities; and, once the scenario has dependency-aware scheduling turned on, link, unlink, or relag dependencies between activities. Twenty-five tools in total. **Setting a description always replaces the previous one** — unlike a note, which the AI can only add to. An empty description clears it. Ask the AI to set it when an activity is first created, or turn on Read Mode first so it can see what it would replace.

Bulk & import calls	Seven of the tools are bulk tools that act on many items in one call — up to 100 activities, milestones, or notes, or 500 dependencies or milestone assignments — plus one that imports a whole schedule (activities, milestones, assignments, and dependencies together) in a single call. A bulk or import call is all-or-nothing at the point you submit it, but each item inside applies on its own — a duplicate id or an invalid estimate skips just that item, not the whole call; the live feed and a follow-up read show exactly what landed. Bulk-appending notes is the one bulk tool that isn't safe to blindly resend — unlike the others, resending it adds every note again (see BUILDING AND UPDATING IN BULK). An import that includes dependencies is the one exception: if dependency-aware scheduling isn't on for the target scenario, the entire import — activities and milestones included — is declined, not partially applied.
What the AI can never do	The AI can never delete an activity, a milestone, or a checklist/deliverable item once added — only you can. Removing a dependency link is the one exception, since an edge carries no content of its own and is trivial to re-add if it was a mistake. The AI can never turn dependency-aware scheduling on or off for a scenario, in either direction — that decision belongs to you alone, in Settings.
Read Mode (optional)	Shares a computed snapshot of your current schedule with the AI — activities (including whether one already has a description), estimates, computed dates and critical path, milestones, dependencies, and a recommended distribution per activity. Required before asking the AI to link, unlink, or relag a dependency, and before asking it to reorder activities (it must re-read the complete, current activity-id list immediately beforehand). Generally needed whenever the AI works with an activity or milestone it didn't just create itself — without it, the AI has no way to learn an existing id. Leave it off if you only want the AI building fresh, from ids it mints itself — most bulk tools work without Read Mode too, the exceptions being bulk dependency creation, an import that includes dependencies, and reorder. Toggle at any time from the Connect AI session panel. Note on freshness: if you've just run the Monte Carlo simulation, the AI can see whether a result is present, but a present result isn't a guarantee it reflects every recent change — ask it to have you re-run the simulation if you need current numbers.
Session data & expiry	Pairing codes are single-use and expire after 15 minutes. Session data (the ops queue and any Read Mode snapshot) is stored in Firebase and auto-purged after 7 days. Click Disconnect to end a session and delete the snapshot immediately. Full details at spertsuite.com/ai-privacy.
AI-built schedule limits	Each scenario holds up to 500 activities, 2,000 dependencies, and 100 milestones; a project holds up to 50 scenarios. These are the same limits that apply to a schedule you build by hand. A checklist or deliverables list on one activity holds up to 50 items; an appended note or a set description is capped at 2,000 characters. A note that would overflow on append is rejected outright, not silently cut short — ask the AI to keep it concise or split it across activities. A description that's too long is rejected the same way, but since setting one always replaces the previous text rather than appending, ask for a shorter one instead. Per-call bulk limits: up to 100 activities or 100 milestones per bulk-create call; up to 500 dependencies or 500 milestone assignments per call; up to 100 activities per bulk-update or bulk-append-notes call (fewer for the latter if notes carry real content); a schedule import accepts all four sections in one call, each capped the same way as its own bulk tool. A very large ask (say, 300 activities) still needs more than one call — the AI should split it into batches on its own.

Tips & caveats

Enable Read Mode before asking the AI to link dependencies, reorder activities, or to update, annotate, or reference an activity or milestone you already created yourself. Undo (Ctrl+Z) works on every AI change; a run of AI edits with no edit of your own in between collapses into a single undo step, so reverting an AI-built schedule is usually one keystroke, not one per activity — this applies to a bulk call or an import too. Don't switch projects mid-session — Connect AI is scoped to the project open when you clicked Connect. The AI never deletes an activity, milestone, or checklist/deliverable item — only dependency links can be removed by the AI, and only you can turn dependency-aware scheduling on or off. A note the AI appends is added to whatever is already there, never overwritten — ask it to add another note rather than expecting it to edit an existing one. Bulk-appending notes is the one bulk tool where resending a call is not automatically safe — it adds every note again rather than skipping what already landed. A description works the opposite way — setting one replaces whatever was there before, so turn on Read Mode first if you want the AI to see the current text before changing it. **Reordering asks for the complete, current set of activity ids** — if the schedule changed since the AI last read it, the reorder is refused outright (never partially applied) and the AI is told to re-read and rebuild the list. In a sequential scenario, reordering changes start/finish dates; in a dependency-aware scenario it only changes the display order and no dates move — either way the simulation is cleared and needs a re-run.

ESTIMATING & ANNOTATING ACTIVITIES WITH THE AI

Once a project exists, you can ask the AI to help build out the activities themselves — estimates, a probability distribution, and supporting detail. Every activity needs a three-point estimate (minimum, most likely, and maximum — in working days); if you don't specify a probability distribution, the AI lets the app recommend one based on the shape of your estimate, the same way the app would for a person entering it by hand. You can also ask the AI to override that recommendation with a specific distribution (normal, log-normal, triangular, or uniform) if you have a reason to. For work with no real duration uncertainty — a booked training week, a fixed contractual notice period — the AI can set all three points to the same number. That is a valid estimate in every distribution: the activity contributes its length to the schedule without contributing any variance to the simulation. Once activities exist, the AI can annotate them: writing or updating a short plain-language description of what the activity entails, appending a note that explains its reasoning, adding checklist or deliverable items, and marking existing items complete.

Example prompts that work well:

Build activities from a description	"I'm scheduling a kitchen remodel. Create activities for demo, rough plumbing, rough electrical, drywall, cabinets, countertops, and final inspection, with a reasonable three-point estimate for each in working days."
Ask for a distribution recommendation	"For the 'countertop fabrication' activity, what distribution would you recommend given a 5/10/25-day estimate, and why?" (The AI can explain the recommendation without changing anything, or apply it if you confirm.)
Pin a fixed-duration activity	"'Vendor training' is a booked five-day course — set min, most likely, and max all to 5 so it doesn't add any uncertainty to the simulation."
Write a short description	"Write a one-sentence description for 'rough electrical' that a teammate could read six months from now to know what the activity covers."
Record its reasoning	"Add a note to 'rough electrical' explaining that the estimate is wider than usual because we're waiting on a permit inspection that could add a week."
Add a review checklist	"Add a checklist to 'final inspection' with items for permit sign-off, punch-list walkthrough, and client sign-off."
Mark something done	"Mark the 'permit sign-off' checklist item on 'final inspection' complete."

Note on notes: the AI can only add to an activity's notes, never overwrite or remove what's already there. If you want the AI's reasoning kept separate from your own notes, say so — it will still land in the same field, just appended after whatever's already written.

Note on descriptions: unlike a note, setting a description replaces whatever was there before — the AI can't merge it with existing text, and an empty description clears it entirely. Ask the AI to write one when it first creates an activity, or turn on Read Mode so it can see the current description before changing it.

Building or annotating a lot of activities at once? See BUILDING AND UPDATING IN BULK below — the AI can create or update dozens of activities in a single call instead of one at a time, seed a description and an initial note for each one right at creation, or append a note to many existing activities in one call.

BUILDING AND UPDATING IN BULK

For a schedule of any real size, ask the AI to work in bulk instead of one activity at a time — it can **create** up to 100 activities, 100 milestones, or 500 dependencies or milestone assignments in a single call, **update** up to 100 existing activities' names, estimates, confidence, distribution, or descriptions in one call, **append a note** to up to 100 existing activities in one call, or **import a whole schedule** — activities, milestones, assignments, and dependencies together — in one call. A 27-activity, 34-dependency, 5-milestone project that would otherwise take well over a hundred individual tool calls fits in one import.

Example prompts that work well:

Build many activities at once	"Create all fourteen activities for this kitchen remodel in one call, with a three-point estimate and a one-line description for each."
Update many estimates at once	"Widen the estimate range on every activity tagged 'permitting' by about 20%, in one call."
Import a whole schedule	"Here's the full activity list, dependencies, and milestones from our planning doc — build the entire schedule in one call, including the dependencies (dependency-aware scheduling is already on)."
Annotate many activities at once	"Add a note to every activity in this schedule flagging the risks specific to that task — do it in one call rather than one activity at a time."

Note on partial results: a bulk or import call is all-or-nothing only at the point you submit it — once it's queued, each item applies on its own. A duplicate id, an invalid estimate, or a reference to something that doesn't exist skips just that item; everything else in the call still goes through. The live activity feed (and a follow-up read from the AI) shows exactly what applied and what didn't, and why. **One exception to keep in mind:** bulk-appending notes isn't safe to blindly resend the way the other bulk tools are — a repeated call adds every note again, so if one partially fails, ask the AI to resend only the activities that were skipped.

Importing dependencies is the one all-or-nothing exception: it requires Read Mode and a scenario with dependency-aware scheduling already on. If that setting is off — or gets turned off between the AI's call and your browser applying it — the **entire import is declined**, including any activities and milestones in the same call. Nothing is half-built. An import with no dependencies section needs neither Read Mode nor dependency-aware scheduling.

Good to know: in an import, if an activity or milestone is skipped for a reason that means it was never created, anything that depended on it (a milestone assignment, a dependency edge) is skipped too, reported as "not found" rather than silently attaching to nothing. A skip for being a duplicate does not cascade this way, since the entity already exists — which is what makes re-running a completed import safe.

LINKING ACTIVITIES WITH DEPENDENCIES

If your activities have real dependencies — one can't start until another finishes — you can ask the AI to link them. This requires two things you set up yourself first: **Read Mode** must be on, and the scenario must already have **dependency-aware scheduling** turned on in Settings. The AI can explain what turning that setting on would mean for your schedule, but it can never flip it for you, in either direction — if you ask it to link activities before the setting is on, it will tell you to enable it yourself rather than silently creating links that don't yet affect anything. Once both are in place, the AI can create, remove, or adjust the lag and type of a dependency between two activities.

Example prompts that work well:

Turn on dependency mode first	"Should this schedule use dependency-aware scheduling? Explain what that means before I decide." (Once you've turned it on yourself in Settings, tell the AI it's ready.)
Link two activities	""Countertops' can't start until 'rough plumbing' and 'rough electrical' are both done — add those dependencies."
Adjust lag or type on an existing link	"Change the lag between 'drywall' and 'cabinets' to 3 working days, since paint needs to cure first."
Remove a link that's wrong	"Actually, 'cabinets' doesn't depend on 'countertops' — remove that dependency."
Link many dependencies at once	"Here are all 34 dependencies for this schedule — create them all in one call." (Same Read Mode + dependency-aware scheduling requirement as a single link.)

Note on cycles: the AI can't create a dependency that would form a loop (activity A depending on B depending back on A) — the tool refuses the specific link that would close the loop, and tells you why, rather than corrupting the schedule. In a bulk call, an acyclic set of edges applies in full regardless of what order you listed them in; if the set (together with edges you already have) contains a cycle, only the edges that would close it are skipped — fix those and resend just the skipped ones. A loop can still reach a scenario by another route — a project JSON file imported from elsewhere. The import preview now names any file carrying a circular chain, a link to an activity that isn't in the file, or an activity depending on itself, and the scenario shows a **Dependency Cycle** banner pointing at the Dependencies panel. Break the loop there before asking the AI to work on that scenario.

Good to know: a scenario's dependency-aware setting is genuinely a one-way door you control, not the AI — if you're not sure whether your schedule needs it, ask the AI to talk through the tradeoff first. Sequential scheduling (activities running one after another in the order they're listed) is often enough for straightforward work; dependency-aware scheduling is worth it once activities can genuinely run in parallel or have real technical ordering constraints.

REORDERING ACTIVITIES

You can ask the AI to put a scenario's activities in a different order — for a client-facing presentation, for example, or to match a phase you're planning around. This requires **Read Mode**: the AI must re-read your project immediately before reordering, since it has to supply the complete, current set of activity ids in the order you want. **What reordering actually changes depends on the scheduling mode:** in a **sequential** scenario, activities run in list order, so reordering changes start and finish dates. In a **dependency-aware** scenario, the schedule follows the dependency graph instead, so reordering changes the **display order only** — no dates move. Either way, the simulation is cleared and needs a re-run, matching what happens when you drag-reorder by hand.

Example prompts that work well:

Reorder for a presentation	"Put the activities in the order I'd present them to the client: foundation work first, then the visible finish work, then inspections last."
Reorder to match a phase plan	"Re-read the project, then reorder the activities to match the three phases we discussed."

Note on a stale list: the AI must supply the exact current set of activity ids — not more, not fewer. If your project changed since the AI last read it (an activity was added, removed, or renamed with a new id), the reorder is refused outright as a **stale order**, never partially applied — the AI is told to re-read the project and rebuild the list. A list with a repeated id is refused as an **invalid order** for the same reason: nothing is guessed or half-applied.

Good to know: section-header bands aren't visible to the AI, but they aren't lost either — each band stays attached to the activity it's anchored to, so it moves along with that activity. A visual grouping may shift as a result; that's expected, not a bug.

TROUBLESHOOTING

Symptom	Likely cause	Fix
AI shows fewer tools than expected (e.g. only a few, or missing dependency tools)	Stale cached manifest. AI clients cache the tool list at connect time.	Fully remove the MCP connector and re-add it. A soft Refresh is often insufficient. The connector lists every tool the shared suite server advertises across all SPERT apps — 51 today , of which 25 are the scheduler_ tools this guide covers. The suite total grows as other SPERT apps add tools, so the check that stays reliable is the second number: after re-adding, confirm the full scheduler_ family is present before starting a session.
ChatGPT behaves unexpectedly: loops, does not follow instructions, or reports few tools	Free-tier model downgrade. Token exhaustion can silently drop ChatGPT to a weaker model.	Start a new conversation. Confirm you are on a capable model tier (GPT-4o or equivalent).
Tool names appear in the AI's reply but nothing changes in the app	The AI has no live MCP connection — it is reading the prompt but cannot reach the server.	The AI assistant must be configured with the SPERT MCP server (see One-Time Setup). Paste the Copy Prompt only after the connector is set up. Tool names in a reply without any actual changes is the classic sign of a missing connection.
appVersion shows an old version number after a deploy	Browser tab is using a cached JS bundle from before the deploy.	Open a new browser tab (do not reuse a tab opened before the deploy). Create a fresh Connect AI session in the new tab.
AI says it linked two activities but the schedule/critical path didn't change	Dependency-aware scheduling is off for this scenario.	The link itself was refused, not silently applied — the AI's own reply should say so. Turn on dependency-aware scheduling for the scenario in Settings yourself, then ask the AI to try again.
AI's dependency tool calls get refused	Read Mode is off, or the scenario doesn't have dependency-aware scheduling enabled.	Turn on Read Mode in the Connect AI panel. If the AI's refusal message mentions the scheduling mode specifically, enable dependency-aware scheduling for that scenario in Settings — the AI cannot do this step for you.
A note the AI wrote got rejected (alone, or in a bulk call)	The appended text would push the note past 2,000 characters.	Ask the AI to keep the note shorter, or split it across activities. In a bulk call, only that one item is skipped — the rest still applies.

AI overwrote an activity's description I wanted to keep	Setting a description always replaces the previous text — unlike notes, it does not append.	This is by design. Ask the AI to check the current description first (turn on Read Mode) before changing one you already wrote, or press Ctrl+Z to undo the change.
Newly created activity has an unexpected probability distribution	The distribution was auto-recommended from the estimate's shape, not chosen by you.	Ask the AI to explain why that distribution was recommended, or to override it with a specific one (normal, log-normal, triangular, or uniform) if you have a reason to prefer a different shape.
Asked the AI to remove a checklist item and it couldn't	The AI can only add items and mark them complete — never remove one.	This is by design — only you can remove a checklist or deliverable item, in the app itself. Ask the AI to mark it complete instead if that's the actual goal.
A bulk call applied some items but skipped others	Each item in a bulk or import call is validated on its own — a duplicate id or bad estimate skips just that item.	This is by design, not a partial failure. Check the live activity feed or ask the AI to read the project back to you — it will name each skip and why (duplicate, cycle, limit reached, invalid, and so on).
An import was declined entirely — nothing in the call happened	The import included dependencies, but the scenario doesn't have dependency-aware scheduling on.	This is all-or-nothing by design when dependencies are included. Turn on dependency-aware scheduling for the scenario yourself, or ask the AI to resend the import without the dependencies section.
AI edits land, but the scenario shows a "Dependency Cycle" banner and no schedule	The scenario's dependencies form a loop — usually from a project file imported from elsewhere, since the AI can't create one.	Open the Dependencies panel and break the loop; the banner points there. The AI's own edits are unaffected and nothing it did caused this — a schedule simply can't be computed until the cycle is gone. If the project came from a JSON import, the import preview flagged it at the time.
AI says a reorder was refused as "stale current order doesn't match"	The project changed since the AI last read it — an activity was added, removed, or renamed.	Nothing was applied — this is a refusal, not a partial reorder. Ask the AI to re-read the project (<code>scheduler_get_project</code>) and rebuild the full list before trying again.
AI says a reorder was refused for a repeated activity id	The requested order listed the same activity id more than once.	Ask the AI to rebuild the list with each activity id appearing exactly once, in the order you want.